
Dive into Deep Learning

xAI-Proj-B: Bachelor Project Explainable Machine Learning

Florian Gutbier*

Otto-Friedrich University of Bamberg
96047 Bamberg, Germany
XXX@stud.uni-bamberg.de

Marius Ludwig Bachmeier†

Otto-Friedrich University of Bamberg
96047 Bamberg, Germany
XXX@stud.uni-bamberg.de

Andreas Schwab‡

Otto-Friedrich University of Bamberg
96047 Bamberg, Germany
andreas-franz.schwab@stud.uni-bamberg.de

Abstract

//TODO: write abstract? Was muss da rein?

1 Introduction

The emergence of deep learning has led to unprecedented advances in various fields, including medical image analysis. This project seeks to explore the fundamental principles of deep learning and to leverage its potential in a practical setting. Our investigation is divided into two parts: First, we focus on the classification of handwritten digits using the MNIST dataset (Lecun et al., 1998), followed by a more complicated challenge of classifying nine different tissue patterns within the PathMNIST dataset (Kather et al., 2018, 2019), a subset of MedMNIST (Yang et al., 2021). These tasks not only serve as a basis for understanding the mechanisms of deep learning, but also highlight the impact of the application of neural networks in medical diagnostics, emphasizing the relevance of our research.

The initial phase of our project was dedicated to learning the basics using the MNIST dataset, which was chosen for its many resources and tutorials to help us get started with deep learning. Our goal was to establish a basic understanding of the field by examining and comparing the results obtained from tuning various hyperparameters. In this phase, we developed a basic convolutional neural network (SimpleCNN) to introduce us to the architectures of neural networks and their ability to classify different digits.

The transition to the MedMNIST dataset, in particular the PathMNIST subset, represented a significant increase in the complexity of our project. This phase was crucial as it enabled us to apply and refine advanced techniques, such as experimenting with pre-trained models, testing a wide range of optimizers, and exploring different strategies for data pre-processing and augmentation. The PathMNIST subset was chosen to emphasize the critical importance of neural networks in medical image analysis. By tackling the classification of tissue patterns, we have not only explored some technical intricacies of deep learning, but also contributed to an area where such technologies could potentially revolutionize diagnostic methods.

This report is structured by first presenting some theoretical foundations that are crucial for understanding the methods used, including dropout layers, batch normalization and ReLU. Next, we

*Degree: B.Sc. AI, matriculation #: 12345678

†Degree: B.Sc. AI, matriculation #: 12345678

‡Degree: B.Sc. AI, matriculation #: 2017990

explain the architectures of different deep learning models that were investigated in the project, from our initial SimpleCNN to more complex pretrained models such as AlexNet, ResNet and Xception (Chollet, 2017). Subsequently, a brief overview of the used datasets MNIST and MedMNIST is given, which form the basis for a deeper investigation. We then present the results obtained with each architecture, addressing the specific challenges. The discussion section provides a critical evaluation of our results and leads to a reflective conclusion about the lessons learned and the potential impact of our research on medical image analysis.

2 Methods

2.1 Theoretical Foundations

//TODO

2.2 Model Architectures

2.2.1 SimpleCNN

Our SimpleCNN model contains two primary convolutional layers. The choice of this architecture was motivated by the goal of understanding the basic mechanisms of neural networks in processing and classifying image data. The model uses LeakyReLU activation to avoid the vanishing gradient problem and to ensure effective backpropagation even at small gradient values. The first convolution layer uses 32 filters with a kernel size of 3×3 , a stride of 1 and ‘same’ padding, which preserves the dimension of the input images. This is followed by a max-pooling layer with a kernel size of 2, which aims to reduce the spatial dimensions while keeping the important features to optimize the model’s ability to detect patterns without considerable data loss. Subsequently, a second convolution sequence increases the depth to 64 filters, improving the model’s ability to extract more features. This setup is again followed by LeakyReLU activation and max-pooling, further refining the feature extraction process. The architecture concludes with a linear layer that maps the high-level features to the output classes, with a softmax function to interpret the outputs as class probabilities.

As our project progressed and our understanding deepened, we enhanced the initial SimpleCNN model by introducing an additional convolutional layer and incorporating batch normalization and dropout techniques, aiming for improved accuracy and generalization. The modified SimpleCNN architecture begins with a convolutional layer designed for single-channel (grayscale) images to match the format of the MNIST dataset. This layer, consisting of 32 filters with a kernel size of 3×3 and ‘same’ padding, is followed by batch normalization and ReLU activation, which promotes nonlinearity while maintaining normalization across the inputs of the network. A dropout rate of 0.25 after max pooling aims to mitigate overfitting by randomly omitting a portion of the features during the training process. As the model progresses, the second and third convolutional layers increase the filter count to 64 and then 128, respectively, each augmented with batch normalization, ReLU activation, max pooling, and a consistent dropout rate of 0.25. This architectural depth ensures the extraction of increasingly complex features essential for recognizing diverse patterns in handwritten digits. The concluding segment of the model comprises a fully connected layer transitioning from the convolutional output to 256 units, followed by batch normalization and a higher dropout rate of 0.5, further combating overfitting. The final linear layer maps these processed features to the ten class outputs corresponding to the MNIST dataset’s digit categories.

The transition from the SimpleCNN model, which was tailored for the MNIST dataset, to the model customized for the PathMNIST challenge required some adjustments to account the different characteristics of the two datasets. The most important change was to adapt the input layer to include 3-channel RGB images for PathMNIST. This change is important to take advantage of the color information that is critical in medical imaging for identifying different tissue types. Despite this adaptation, the core architecture - consisting of convolutional layers, batch normalization, ReLU activation, and dropout layers - remained consistent between models.

For a detailed implementation of all versions of the model, refer to the Appendix Section A.1.

2.2.2 ResNet

//TODO

Resnet18 //TODO

Resnet50 //TODO

ResnetXX //TODO

2.2.3 Xception

To understand the Xception model, one must become familiar with depth-separable convolutions. This technique is a convolution operation that divides into two different stages to increase computational efficiency. Initially, a depth-wise convolution applies a single filter to each input channel. Subsequently, a point-wise convolution - characterized by a 1×1 kernel - combines the outputs from the depth-wise step across the channels. This factorization significantly reduces the number of parameters and calculations and enables more efficient training. Batch normalization follows each convolution and promotes stable learning by normalizing the ReLU activations of the layer.

As seen in Figure 6 in the Appendix the architecture of Xception, as detailed by (Chollet, 2017), is structured into three primary flows. The entry flow prepares the network with initial convolutions and pooling to create condensed feature maps from the input images. It starts with two standard convolutions, followed by a series of separable convolutions that increase the depth while compressing spatial dimensions. This is achieved using a combination of 1×1 convolutions for channel-wise feature processing and 3×3 convolutions for capturing spatial information, each followed by max pooling to halve the feature map dimensions progressively. This step reduces the initial input image size from $299 \times 299 \times 3$ to $19 \times 19 \times 728$. Central to Xception’s design is the middle flow, which repeatedly applies depthwise separable convolutions to process and refine the features. This part of the network is designed to be repeated eight times, allowing the model to learn increasingly complex patterns without affecting the dimensions. The exit flow then expands the feature maps through additional separable convolutions, incorporating a mix of channel-wise and spatial feature extraction before concluding with global average pooling. This step reduces each map to a single vector, capturing the essence of the input data in a form suitable for classification. The network concludes with a logistic regression layer, such as softmax, to output the final class probabilities.

3 Experiments

Our experimental framework focuses on the use of two central datasets: MNIST and MedMNIST, each of which is briefly introduced here. This section also describes the experiments performed with different model architectures and presents the resulting findings.

3.1 MNIST

The “Modified National Institute of Standards and Technology” dataset (Lecun et al., 1998) comprises a collection of 70,000 handwritten digits carefully divided into a training set of 60,000 images and a test set of 10,000 images. Each digit is represented in a grayscale image of 28×28 pixels, and offers a wide range of styles and shapes. This dataset is widely recognized for its simplicity and effectiveness in benchmarking classification algorithms, making it an ideal starting point for those new to deep learning. Some examples of the dataset can be seen in Figure 1



Figure 1: Example grayscale images of the MNIST dataset

3.2 MedMNIST

MedMNIST, a more specialized and challenging dataset than MNIST, is tailored for medical image classification tasks. It extends the concept of handwritten digit classification to a diverse range of medical imaging modalities, including dermatology or radiology. Unlike MNIST’s uniform format,

MedMNIST encompasses 12 subsets for 2D and 6 subsets for 3D data. For our project, we focused on the PathMNIST (Kather et al., 2018, 2019) subset, which includes “100,000 non-overlapping image patches from hematoxylin and eosin stained histological images, and a test dataset [...] of 7,180 image patches from a different clinical center” (Yang et al., 2021). The images could be classified into nine different types of tissues.

Initially, the images were of high resolution ($3 \times 224 \times 224$ pixels), but the authors of MedMNIST resized them to $3 \times 28 \times 28$ pixels. The 100,000 training images were then originally divided into training and validation sets in a 9:1 ratio. Some examples of the various images can be seen in Figure 2.



Figure 2: Example images of the MedMNIST dataset.

The PathMNIST subset provides a unique challenge by introducing the complexity of medical image analysis. It requires the use of advanced deep learning techniques and models to accurately classify different types of tissue, making it an excellent progression from the simpler MNIST dataset.

3.3 SimpleCNN

After establishing the basic architecture of our SimpleCNN (as described in Section 2.2.1), our next goal was to effectively classify the first digits of the MNIST dataset. This required the selection of initial hyperparameters to begin the training and evaluation process. Without prior benchmarks, we initially decided on a batch size of 8, set the number of epochs to 5 and a learning rate of 0.1. However, this initial configuration led to less satisfactory results, which made us to revise our hyperparameters to improve performance.

The training process was completed relatively quickly so that we could implement a function to test the hyperparameters. This required us to set default values for each training and evaluation session, with certain values being adjusted at each iteration. For example, we set the batch size to 32, the learning rate to 0.001 and the epochs to 15. In subsequent runs, we changed one hyperparameter at a time, while keeping the others constant. This means that we ran training sessions with learning rates of 0.001, 0.01 and 0.1, respectively, while keeping the default values for batch size and epochs. Similarly, we experimented with different batch sizes (16, 32, and 64) and varied the number of epochs (5, 10, 15, and 20), to further understand their individual contributions to the model’s accuracy and training efficiency.

After analyzing the experimental data (see Appendix, Table 1), we discovered the most effective hyperparameters for our SimpleCNN model, and interestingly, we found that the default values were already optimal. We obtained these results with the original SimpleCNN architecture. Subsequently, we applied these hyperparameters to train an updated version of the SimpleCNN. To assess the impact of the changes, we compared the performance of the original and updated architectures. The comparative analysis is shown in Figure 3.

For the second challenge, which focused on the PathMNIST dataset, our approach exclusively used the adapted version of our SimpleCNN architecture due to its better performance. To combat overfitting, we introduced an early stopping mechanism. The core of this experiment was to assess the impact of incorporating batch normalization and dropout techniques. We conducted three separate trials: one with both batch normalization and dropout, one with only dropout, and another with only batch normalization. The outcomes of these trials can be seen in Figure 4.

3.4 AlexNet

//TODO




Figure 3: Comparison of validation accuracy across epochs between the initial and the adapted SimpleCNN model, highlighting the overall improved performance of the adapted model.

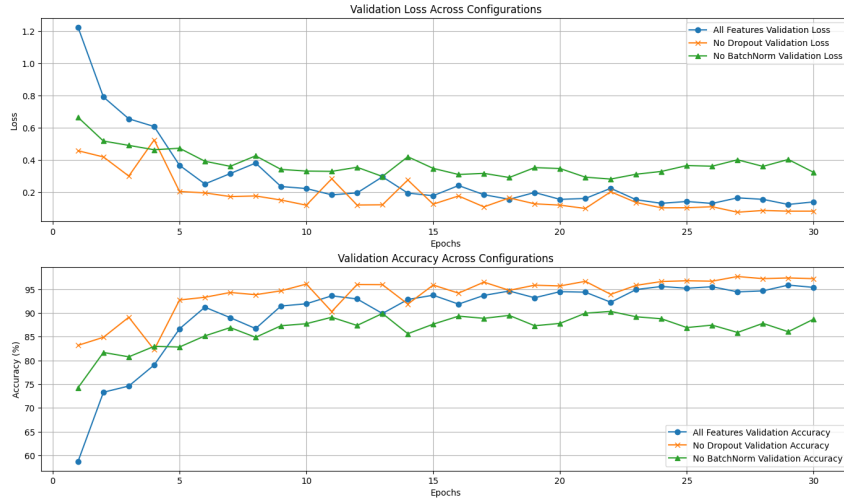


Figure 4: Performance comparison of SimpleCNN variations on the PathMNIST dataset over 30 epochs.

3.5 ResNeT

//TODO

3.6 Xception

The implementation of the Xception architecture presented us with particular challenges right from the start, mainly because Xception is not included in the PyTorch model library. We used the “pretrainedmodels” library (Cadene, 2018) to gain access to a pre-trained Xception model, although this is outdated (last update in April 2020). While the library’s pre-trained Xception model was operational, its use led to some further complications. One major hurdle was adjusting the input size to the Xception default requirement of 299×299 pixels, which is a significant increase from PathMNIST’s original 28×28 pixels. Attempting to upscale images to this size proved to be computationally intensive and resulted in crashes, both on local PCs and when using Google’s V100 GPUs in Google Colab (Google, 2024). To mitigate this problem, we initially resized the images to a more manageable 168×168 pixels. This adjustment allowed us to maintain the batch size of 64, which is consistent with our previous models for PathMNIST. However, the training time for a single epoch with a batch size of 64 was still very long - about 16 minutes on a V100 GPU via Colab and about 4 hours on local machines. Given the unreliability of Colab due to possible interruptions when GPUs are not available, as well as the cost of using a V100 GPU, we further reduced the stack size to

32. This adjustment reduced the epoch training time to about 10 minutes and allowed for a smoother local training process.

For training the Xception model, we used the AdamW optimizer (Loshchilov and Hutter, 2019), known for its effectiveness in handling weight decay and improving generalization by modifying the traditional Adam optimizer’s handling of the learning rate.

With these changes, we were able to successfully run our experiments with the Xception model. The performance results are shown in Figure 5:




Figure 5:

4 Discussion

In our initial experiments with the MNIST dataset using the SimpleCNN model, we compared the effectiveness of the ReLU and LeakyReLU activation functions. We found no significant difference in performance between them, which led us to continue with ReLU due to the simplicity of our model. For the PathMNIST challenge, our experiments with SimpleCNN showed that using only batch normalisation yielded slightly better results than combining it with dropout layers, suggesting that the latter might introduce unnecessary complexity that reduces model performance in this particular task.

Using the Xception architecture for the PathMNIST dataset significantly improved the results, demonstrating the value of deeper, more complex models for processing complicated datasets. However, this came at the expense of computational resources, emphasising the need to strike a balance between the complexity of the model and the practicalities of computational effort. This experience highlighted the critical trade-offs in deep learning between achieving high accuracy and coping with resource constraints, particularly in the context of complex medical imaging tasks.

5 Conclusion

//TODO

References

- R. Cadene. Pretrained models for pytorch. <https://github.com/Cadene/pretrained-models.pytorch>, 2018. Accessed: 14.01.2024, PyPI link: <https://pypi.org/project/pretrainedmodels>.
- F. Chollet. Xception: Deep learning with depthwise separable convolutions, 2017.
- Google. Google colaboratory. <https://colab.research.google.com>, 2024. Accessed: 14.01.2024.

- J. N. Kather, N. Halama, and A. Marx. 100.000 histological images of human colorectal cancer and healthy tissue, 2018. URL <https://doi.org/10.5281/zenodo.1214456>.
- J. N. Kather, J. Krisam, P. Charoentong, T. Luedde, E. Herpel, C.-A. Weis, T. Gaiser, A. Marx, N. A. Valous, D. Ferber, et al. Predicting survival from colorectal cancer histology slides using deep learning: A retrospective multicenter study. *PLoS Medicine*, 16(1):e1002730, 2019.
- Y. Lecun, L. Bottou, Y. Bengio, and P. Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998. doi: 10.1109/5.726791.
- I. Loshchilov and F. Hutter. Decoupled weight decay regularization, 2019.
- J. Yang, R. Shi, and B. Ni. Medmnist classification decathlon: A lightweight automl benchmark for medical image analysis. In *IEEE 18th International Symposium on Biomedical Imaging (ISBI)*, pages 191–195, 2021.

Declaration of Authorship

Ich erkläre hiermit gemäß § 9 Abs. 12 APO, dass ich die vorstehende Projektarbeit selbstständig verfasst und keine anderen als die angegebenen Quellen und Hilfsmittel benutzt habe. Des Weiteren erkläre ich, dass die digitale Fassung der gedruckten Ausfertigung der Projektarbeit ausnahmslos in Inhalt und Wortlaut entspricht und zur Kenntnis genommen wurde, dass diese digitale Fassung einer durch Software unterstützten, anonymisierten Prüfung auf Plagiate unterzogen werden kann.

Bamberg, February 28, 2024

(Place, Date)

(Signature)

Bamberg, February 28, 2024

(Place, Date)

(Signature)

Bamberg, February 28, 2024

(Place, Date)

(Signature)

A Appendix

A.1 SimpleCNN architectures

Initial version of our SimpleCNN, including two convolutional layers:

```
1 class SimpleCNN(nn.Module):
2     def __init__(self, num_classes=10):
3         super(SimpleCNN, self).__init__()
4         self.conv1 = nn.Sequential(
5             nn.Conv2d(
6                 in_channels = 1,
7                 out_channels = 32,
8                 kernel_size = 3,
9                 stride=1,
10                padding="same"
11            ),
12            nn.LeakyReLU(),
13            nn.MaxPool2d(kernel_size=2),
14        )
15        self.conv2 = nn.Sequential(
16            nn.Conv2d(32, 64, 3, 1, "same"),
17            nn.LeakyReLU(),
18            nn.MaxPool2d(kernel_size=2),
19        )
20        self.out = nn.Linear(64*7*7, num_classes)
21
22    def forward(self, x):
23        x = self.conv1(x)
24        x = self.conv2(x)
25        x = x.view(-1, 64*7*7)
26        output = self.out(x)
27        return torch.log_softmax(output, dim=1)
```

Structure of the improved version of the SimpleCNN using three convolutional layers, Batch normalization and Dropout:

```
1 class SimpleCNN(nn.Module):
2     def __init__(self, num_classes=10):
3         super(SimpleCNN, self).__init__()
4         self.conv1 = nn.Sequential(
5             nn.Conv2d(1, 32, kernel_size=3, stride=1, padding="same"),
6             nn.BatchNorm2d(32),
7             nn.ReLU(),
8             nn.MaxPool2d(kernel_size=2),
9             nn.Dropout(0.25)
10        )
11        self.conv2 = nn.Sequential(
12            nn.Conv2d(32, 64, kernel_size=3, stride=1, padding="same"),
13            nn.BatchNorm2d(64),
14            nn.ReLU(),
15            nn.MaxPool2d(2),
16            nn.Dropout(0.25)
17        )
18        self.conv3 = nn.Sequential(
19            nn.Conv2d(64, 128, kernel_size=3, stride=1, padding="same"),
20            nn.BatchNorm2d(128),
21            nn.ReLU(),
22            nn.MaxPool2d(2),
23            nn.Dropout(0.25)
24        )
25        self.fc1 = nn.Linear(128 * 3 * 3, 256)
26        self.fc_bn = nn.BatchNorm1d(256)
27        self.dropout_fc = nn.Dropout(0.5)
28        self.fc2 = nn.Linear(256, num_classes)
29
30    def forward(self, x):
31        x = self.conv1(x)
32        x = self.conv2(x)
33        x = self.conv3(x)
34        x = x.view(-1, 128 * 3 * 3)
35        x = F.relu(self.fc_bn(self.fc1(x)))
36        x = self.dropout_fc(x)
37        x = self.fc2(x)
38        return torch.log_softmax(x, dim=1)
```

Structure of the improved version of the SimpleCNN for the PathMNIST dataset:

```

1  class SimpleCNN(nn.Module):
2      def __init__(self, num_classes=10):
3          super(SimpleCNN, self).__init__()
4          self.conv1 = nn.Sequential(
5              nn.Conv2d(3, 32, kernel_size=3, stride=1, padding="same"),
6              nn.BatchNorm2d(32),
7              nn.ReLU(),
8              nn.MaxPool2d(kernel_size=2),
9              nn.Dropout(0.25)
10         )
11         self.conv2 = nn.Sequential(
12             nn.Conv2d(32, 64, kernel_size=3, stride=1, padding="same"),
13             nn.BatchNorm2d(64),
14             nn.ReLU(),
15             nn.MaxPool2d(2),
16             nn.Dropout(0.25)
17         )
18         self.conv3 = nn.Sequential(
19             nn.Conv2d(64, 128, kernel_size=3, stride=1, padding="same"),
20             nn.BatchNorm2d(128),
21             nn.ReLU(),
22             nn.MaxPool2d(2),
23             nn.Dropout(0.25)
24         )
25         self.fc1 = nn.Linear(128 * 3 * 3, 256)
26         self.fc_bn = nn.BatchNorm1d(256)
27         self.dropout_fc = nn.Dropout(0.5)
28         self.fc2 = nn.Linear(256, num_classes)
29
30     def forward(self, x):
31         x = self.conv1(x)
32         x = self.conv2(x)
33         x = self.conv3(x)
34         x = x.view(-1, 128 * 3 * 3)
35         x = F.relu(self.fc_bn(self.fc1(x)))
36         x = self.dropout_fc(x)
37         x = self.fc2(x)
38         return torch.log_softmax(x, dim=1)

```

A.2 Pretrained Model Architectures

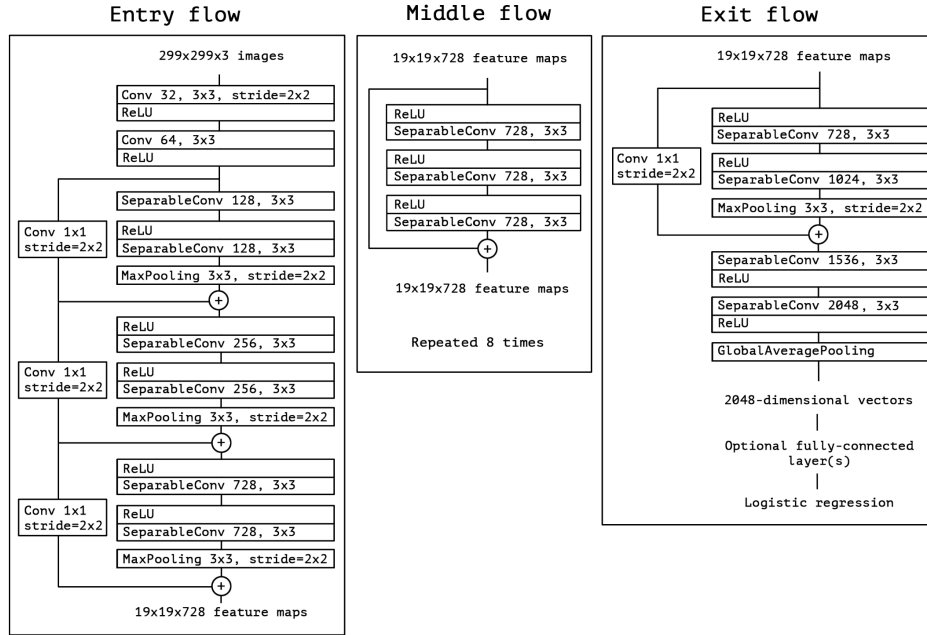


Figure 6: Xceptions architecture as outlined by Chollet.

A.3 Additional Tables

Table 1: Impact of Learning Rate, Batch Size, and Epochs on Accuracy in order to determine the best set of Hyperparameters for the SimpleCNN

Learning Rate		
Learning Rate	Average Train Accuracy (%)	Average Validation Accuracy (%)
0.001	99.6	98.9
0.01	97.7	97.2
0.1	10.3	10.0

Batch Size		
Batch Size	Average Train Accuracy (%)	Average Validation Accuracy (%)
16	99.7	98.9
32	99.5	98.9
64	99.4	98.9

Epochs		
Epochs	Average Train Accuracy (%)	Average Validation Accuracy (%)
5	99.4	98.8
10	99.6	99.0
15	99.6	99.0
20	99.7	99.0